A Review of Genome Graph Tools

Njagi Mwaniki and you if you wanna ;)

June 3, 2020

Abstract

Contents

List of Figures

1 Introduction

There are four major overlapping steps that need to be carried out after DNA sequencing to get actionable results. 1) fragment assembly which involves recreating the original genetic sequence; 2) alignment, mapping, indexing and querying which are part of the assembly stage but can be done later to query the relatedness of sequences; and 3) Visualization to see these changes.(Flicek & Birney, 2009)

expand intro

2 Fragment Assembly

The fragment assembly problem, reconstructing long contigs from reads (Chikhi, Limasset, & Medvedev, 2016), started with Sanger sequencing. Like solving a puzzle of sequences the assembly problem can be framed as finding a set of common super-strings from a set of substrings to which tools take either a Eulerian, connect the nodes, or a Hamiltonian, connect the edges, approach.

The first solution was SEQAID, short for **seq**uencing **aid**, (Peltola, Söderlund, & Ukkonen, 1984) which used the Overlap Layout Consensus (OLC) framework, an intuitionistic approach to the assembly problem that works in 3 main steps: 1) Overlap: find all overlaps among the reads, 2) Layout: assemble a graph of all the reads based on their overlaps, and 3) Consensus: infer the consensus sequence.

What it does is find the overlaps between the given set of reads and then lays them out to form a DAG. The graph, however, has multiple paths that are the genome or superstring in question. Which is where the consensus comes in. It used ... to infer the consensus.
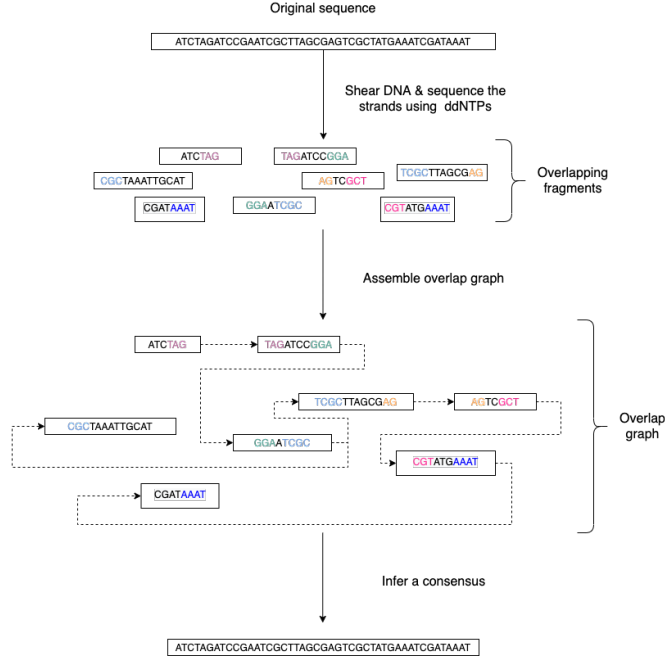
Figure 1: OLC...

With the advent of Second Generation Sequencing (SGS) technologies Roche/454 (www.454.com), Illumina/solexa (www.illumina.com), and AB/Solid (www.appliedbiosystems.com) there was a significant reduction in the cost per sequenced nucleotide (Z. Li et al., 2012) leading to an explosion in the number of reads produced and thus making OLC unfeasible.

OLC based methods were progressively replaced by de Bruijn Graph (DBG) (Idury & Waterman, 1995) based methods whose first tool was EU-LER (Pevzner, Tang, & Waterman, 2001).

The DBG transforms the assembly problem into a computationally simpler problem. While the OLC, through the reads graph, takes an non polynomial time Hamiltonian approach; the DBG, using the k-mer graph, goes about it using a much easier to solve Eulerian approach which works in polynomial time. (Z. Li et al., 2012; Pevzner et al., 2001)

The DBG is an anti-intuitionistic solution whose assembly quality improves with an increase in the read length <cn>. It works through 1) chop the reads into much smaller k-mers, 2) Use all the k-mers to assemble a dBG, 3) Infer the genome sequence on the dBG. These tools made it possible to work with the increasing number of reads produced.
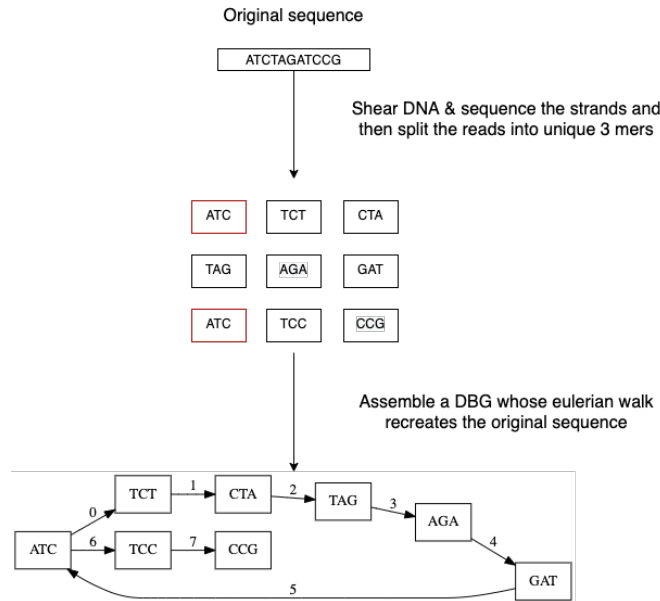
4

Figure 2: DBG...

Lately, Single Molecule Sequencing (long read sequencing) methods have reintroduced OLC for the assembly of long erroneous reads but DBG based methods are still used to correct long reads. The unitig approach involves finding the maximal interval subgraphs in the graph of all read overlaps.

The string graph (Myers, 2005) is a way to infer the genome from the read data without a need for k-mers as is with the de Bruijn graph. It treats the genome as a sequence of repeats and then "encodes" each repeat as a node. Each unique (contiguous) sequence is a node. An Eulerian tour recreates the genome.

Figure 3: String Graph...

3 Alignment and mapping

Alignment involves computing the amount of similarity between two strings also known as the edit distance problem. The solution to the edit distance problem by (Levenshtein, 1966) paved the way for solving the alignment problem.

The first solution was global alignment (Needleman & Wunsch, 1970) in which a sequence query is aligned to the other (reference) in its entirety. It took a dynamic programming approach which worked in square time (had a complexity of $O(n^2)$; it was followed by semi-global alignment (Sellers, 1980) where one sequence (query) is entirely aligned to a substring of the other (reference); then local alignment (Smith & Waterman, 1981), where the alignment can be between any substrings of the two sequences.

In practice when given a set of reads, a complement of each read is generated to be searched against because of the direction of sequencing or inversions. A match can either be exact, matching the pattern exactly, or fuzzy, where a section or all of the string matches the pattern approximately, with minimum edit distance.

With graphs, reads are mapped to paths in the graph instead of linear sequences. Alignment problems grow with the input size (Durbin, 2014) making it hard to align sequences to graphs because of the increased amount of data involved The complexity of an alignment problem is a function of the number of vertices |V| and edges |E| <cn>. In some way you can think of it as mapping to multiple linear sequences that may or may not loop.

4 Indexing

Indexing is a solution to the problem of search given limited computing resources. An index is useful to speed up alignment and make it pragmatic within the given time and memory requirements. It involves reducing the search space so as to reduce the time taken and memory consumed when performing a search. In linear references commonly used indexing approaches are the FM index todo{list tools} whose complexity is O(NM) where there are N variable sites and M sequences (Durbin, 2014). As in alignment, the problem grows even larger with the proliferation of paths in graphs. For graphs, indices like the FM-index backed by the BWT fail to hold<cn> and there's the need for improvements such as that seen in gBWT used in seqwish allowing it to be orders of magnitude faster than VG.

An index can either be static or dynamic. A static index is serialized and saved to disk while a dynamic index is created at runtime and held in memory. Dynamic indices are good with small datasets that change rapidly such as in the construction of a DBG making it suitable for fragment assembly. Static indices are suited for larger datasets that we want to go back to such as a reference genome graph.

Below are some of the approaches taken to solve the problem of indexing

4.1 Burrows-Wheeler Transform

The Burrows-Wheeler Transform (BWT) was introduced by (Burrows & Wheeler, 1994) for string data compression and to this day forms the basis of the bzip compression algorithm.

It works as a preliminary step in the building of indices and also the compression of _____   expound

4.2 Suffix Array

Suffix arrays were introduced by (Manber & Myers, 1990) are arrays of the positions of all the sorted suffixes of a string. A suffix array is a simple, space efficient (stores n integers where n is the length of the string) alternative to the suffix tree <citation needed> whose space requirements are... based on BWT have been used for fast search algorithms

Improvement to the suffix array: (H. Li, 2016) gave the first in-place O(n) time suffix array construction algorithm that is optimal both in time and space, where in-place means that the algorithm only needs O(1) additional space beyond the input string and the output suffix array.

Tools using the suffix array include Bowtie (Langmead, Trapnell, Pop, & Salzberg, 2009), BWA (H. Li & Durbin, 2009), and SOAP2 (R. Li et al., 2009).

4.3   FM Index

Short for Full-text index in Minute space; the FM-index created by (Ferragina & Manzini, 2000) is a full text substring index based on the BWT. It allows compression of the input text while permitting fast substring queries. It can be used to efficiently find the number of occurrences of a pattern within the compressed text, as well as locate the position of each occurrence.

4.4   Positional Burrows-Wheeler Transform

Introduced by (Durbin, 2014) Positional Burrows Wheeler Transform is an algorithm with complexity O(NM) where M sequences and N bi-allelic sites. It derives a representation of the data based on a positional prefix array; an array that holds positions of a given array/set of haplotypes in a larger haplotype array. This prefix array orders them in reverse (ascending) order of their prefixes allowing similar sequences to cluster together.

<Add PBWT table and graphic>

4.5   GBWT/gPBWT

First described (Novak, Garrison, & Paten, 2017) but used in a tool (Sirén, Garrison, Novak, Paten, & Durbin, 2020) it's a compressible representation of a set of haplotypes held in the graph. This allows for efficient match queries in sections of the haplotypes (local alignment). Because of the previously mentioned nature of the positional suffix array to bring together (fairly) similar haplotypes. GBWT lets us have an efficient way of counting the number of haplotypes containing a given sequence.

4.6   Bloom filters

The bloom filter is a probabilistic data structure that can give false positive but never a false negative. It works by hashing data and stores the hash in an array... It is suited for the fragment assembly using DBGs because of its constant time access (Chikhi & Rizk, 2013). It however suffers from poor data localization which led to the use of Blocked Bloom Filters (BBF) [expound] (Putze, Sanders, & Singler, 2010) used in Bifrost (Holley & Melsted, 2019).

4.7 Minimizers

The work of a minimizer is to reduce the search space. It does this by generating kmers from a read and sorting them alphabetically. The k-mer at the top is the minimizer for that read... then binning the result. When a query is made it's prefix is checked against the bin and the rest of the data ignored <is this even accurate?> We can get a minimizer by BBF blocked bloom filter Minimizers (Grabowski, Deorowicz, & Roguski, 2015; Roberts, Hayes, Hunt, Mount, & Yorke, 2004)

4.8 Hash tables

Hash tables involve breaking down the reads into k-mers and storing the kmers into hash tables that point to the original data. When queries are made they're similarly broken down into k-mers of the expected size<citation needed>. Hash based methods when well tuned can be faster than suffix array based methods, because the basic operations are simpler, but they typically require greater memory, particularly in cases where the suffix representation can be compressed as it can be here (Durbin 2014). Many times tools take a hybrid approach; incorporating different aspects of different indexing schemes such as in Minimap (H. Li, Feng, & Chu, 2020).

ensure this citation checks out

5 Genome Graph Tools

The Berkeley Open Assembler (Myers, 2005) uses the string graph 2, a way to infer the genome from the read data without a need for k-mers, and borrows from the unitig algorithm. It treats the genome as a sequence of repeats and then "encodes" each repeat as a node. Each unique (contiguous) sequence is a node. An Eularian tour recreates the genome. Though the original DBG approach does much better than OLC it still has a high memory footprint<citation needed>.

Minia (Chikhi & Rizk, 2013) proposed the use (encoding) of a DBG as a bloom filter (BF) instead of storing the graph in a "traditional" set series of nodes and edges stored in <mention tool>. By design, a bloom filter can give a false positive result but never a false negative therefore the name probabilistic de Bruijn graph. It is obtained by inserting all the nodes of a de Bruijn graph (i.e all k-mers) in a BF. A BF has a search/access time of $O(1)$. They then had an additional structure to remove critical false positives. It showed that the graph can be encoded with as little as 4 bits per node. Drawbacks include 1) The Bloom filter introduces false nodes

and false branching, 2) The global structure of the graph is approximately preserved up to a certain false positive rate.

Bcalm2 (Chikhi et al., 2016) tried to improve this by use of compacted DBG (cdBG). It allowed the problem to be doable on a PC.

<add compaction diagram>

The use of the de Bruijn graph in fragment assembly consists of a multi-step pipeline. The most data intensive steps are usually the first three: nodes enumeration/k-mer counting: the set of distinct k-mers is extracted from the reads Compaction: all unitigs (paths with all but the first vertex having in-degree 1 and all but the last vertex having out-degree 1) are compacted into a single vertex graph cleaning: artifacts due to sequencing errors and poly-morphism are removed from the graph

Minimap (H. Li, 2016) introduced two tools minimap, a raw read over-lapper, and miniasm (H. Li, 2016), an assembler. Minimap uses minimizer sketches, stores k-mers in a hash table, uses sorting extensively. SPAdes also a toolkit does...

Variation graphs embed the paths in the graph. These paths can be used to represent haplotypes. vg, HashGraph, odgi and PackedGraph are dynamic (allow for updates to the graph while xg isn't).

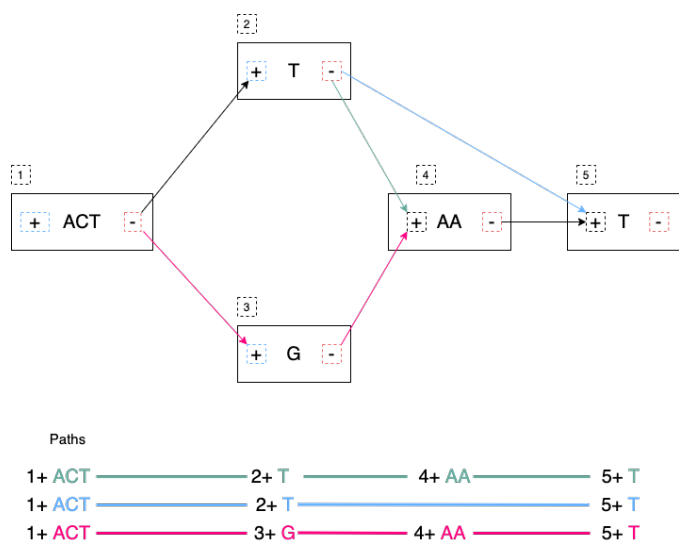|							
Individual 1	A	C	T	G	A	A	T
Individual 2	A	C	T	T	-	-	T
Individual 3	A	C	T	T	A	A	T
Consensus	A	C	T	T	A	A	T

Figure 4: no structural variation

In case of addition of individual 4 containing an inversion at node 2,3 and 4

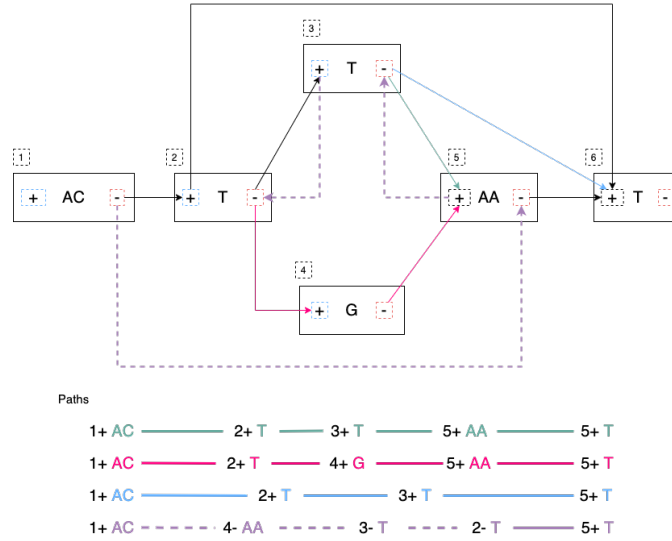|							
Individual 1	A	C	T	G	A	A	T
Individual 2	A	C	T	T	-	-	T
Individual 3	A	C	T	T	A	A	T
Individual 4	A	C	A	A	T	T	T
Consensus	A	C	T	T	A	A	T

Figure 5: contains struct variation

vg (Garrison et al., 2018) used the protobuf library as the graph implementation but was refactored to use the HandleGraph API as of 1.22.0. It is an end to end pangenome graph solution for de novo and incremental graph building but has large memory requirements when it comes to indexing. How it deals with cycles in the graph: unrolls the graph with cycles. The graph holds nodes in a vector. It uses hash tables to map between nodes and ids in a vector that holds the nodes. Paths are stored in a set of linked lists. A hash table maps between nodes and paths. Suffered from a problem of data duplication. Queries involve hash table lookups. xg (Garrison, 2019) a memory-efficient succinct representation of the graph (compared to vg). It being static and having indexing of the paths using positional indices, GBWT, allows it to have fast queries making it good at read mapping and variant calling.

Bluntification (Garg et al., 2018) removing all overlaps between nodes seqwish (`https://github.com/ekg/seqwish`) transforms a set of sequences and alignments (in GFA) into its equivalent variation graph. The large memory requirements of vg are solved through the use of gBWT backed by a Generalized Compressed Suffix Array. It's still a static index odgi (libhandlegraph paper) Optimized Dynamic Graph Interface, uses a dynamic index and uses an in memory variation graph to perform sorting, pruning, transformation, and visualization.

HashGraph (libhandlegraph paper) has speed as its primary goal. Represents a graph as a high performance hash table. Paths are embedded as double linked lists. Edges are in vectors attached to each node they connect. Use an adjacency list which is appropriate for sparse graphs. Appropriate for small graphs <such as viruses> because it trades memory for time. Odgi (libhandlegraph paper) is based on a node centric encoding of the graph that is designed to improve cache coherency when traversing or modifying the graph. It tries to be a pragmatic tool that achieves balance between memory usage and performance. Each nodes seq and edges are encoded in a byte array using a variable length integer, edges are described in terms of the relative offset of a node in a sorted graph. PackedGraph (libhandlegraph paper) is designed to have a low memory footprint. It does this by encoding the graph mainly using linked lists. Baum (Wang, Wang, Li, & Li, 2018) By Adaptive Unique Mapping (BAUM) improved on the OLC framework to improve genome assembly based on SGS paired-end/mate-pair libraries.

BAUM has two modules: construction of the genome unique regions that are taken as the initial contigs iterative assembly, in which scaffolds are built, and contigs are extended and merged, aiming to reconstruct the repetitive regions along the iterations. In this scheme, the repetitive regions are separated by the unique regions Bifrost (Holley & Melsted, 2019) improved on the cdBG by adding colours and takes advantage of concurrency (parallell) to the nodes to keep track of the source of each vertex. Size of colours can grow beyond that of the nodes In the output it stores these colours in a different on a different .bfg$_{colors}$ file. K-mers contained in the unitigs are mapped to their colors representing the input sources (color is represented by an integer from 1 to |C| where C is the number of colors. Colors are stored in a separate array of color containers, each color container is indexed by MPHF (Minimal Perfect Hash Function) library BBHash (Limasset, Rizk, Chikhi, & Peterlongo, 2017). BF has poor data localization because one element is scattered all over which leads to CPU cache misses when inserting and querying are addressed here (Putze et al., n.d.) for this they used (BBF) blocked bloom filter Minimizers (Roberts et al., 2004; Grabowski et al., 2015). BBF works

by building an approximation of the dBG using BBFs to filter our seq errors. BBF containing k-mers is used to build the cdBG.

GraphAlighner (Rautiainen, Mäkinen, & Marschall, 2019) is a tool for aligning long error prone reads to genome graphs. It performs base alignment. It uses (generalizes two linear sequence-to-sequence algorithms to graphs) two strategies: the Shift-And algorithm for exact matching (exact match of a substring to a string) and Myer's bit-vector algorithm for semi-global alignment Aligns sequences to graphs while exploiting bit parallelism. Makes use of Nondeterministic Finite Automaton (NFA). Store an NFA state bitvector for each node and update until no more change is necessary Myer's bit-vector algorithm studies the semi-global sequence-to-graph alignment problem. It seeks to find a path in a directed, node-labelled graph that has the minimum edit distance to the query sequence. Myers' bit-vector alignment algorithm (Myers, 1999) to graphs, which proceeds along the same lines as the Shift-And algorithm, but requires some further algorithmic insights to handle nodes with an in-degree greater than one. Bitvector algo complexity grows approximately linearly with the number of vertices in the graph. The bitvector it uses is the size of the pattern we are searching for. Semi-global alignment is solved through generalizing DP edit distance problem for graphs. Semi-global alignment is used to align a shorter seq against a longer one, reference. Shift-And algorithms (Baeza-Yates and Gonnet, 1992; Domolki, 1964, 1968) performs exact string matching to graphs. Their aim is to find a path in a directed, node-labeled graph that has a minimum edit distance (Levenshtein, 1966) to the query sequence. Shift-And algo finds exact matches between a pattern string and a text string by simulating a nondeterministic finite automaton (NFA) that matches the pattern and then feeding the text to it. Keep shifting the bit-vector by one and bitwise AND-ing the state. Somewhat analogous to exact matching using a window of the size of the pattern. It can handle DAGs and graphs that may contain cycles. For DAGs, process the nodes in topological order (topological sort). For cyclic graphs no sorting.

Minigraph (H. Li et al., 2020), for incrementally constructing reference variation graphs, is a sequence to graph mapper that incrementally constructs a pangenome graph. A graph-based data model and associated formats to represent multiple genomes while preserving the coordinate of the linear reference genome. A straightforward way to represent a pangenome store unaligned genomes in a full-text index that compresses redundancies in sequences identical between individuals (Boucher et al., 2019; Liu, Zhu, et al., 2016; Mäkinen et al., 2010) The other class of methods encodes multiple genomes into a sequence graph, usually by collapsing identical or similar se-

quences between genomes onto a single representative sequence. The results in a pangenome graph.

GraphAlighner performs base alignment but minigraph does not. Minigraph is faster than GraphAligner and uses less memory. Minigraph is more accurate than GraphAligner. This is counter-intuitive given that GraphAligner does base alignment. Close inspection reveals that most mismapped reads by minigraph are mapped to the correct genomic loci but wrong graph paths. On the contrary, most mismapped reads by GraphAligner are mapped to wrong genomic loci. Vg didn't work with their PacBio data. vg allows different regions in one chromosome collapsed to one segment. We call such a graph a collapsed graph. rGFA cannot encode a collapsed graph.

vg-flow (Baaijens, Stougie, & Schönhuth, 2020) attempts to reconstruct all individual haplotypes from a mixed sample at the strain level and to provide abundance estimates for the strains. It does this by...

6   Interfaces and APIs

The field of genome graphs is growing quickly as evidenced by the ever-growing number of tools creating the need for a common way for these tools interact with the data they operate on.

One such solution is libhandlegraph, a declarative approach towards graphs that defines an interface between which tools interact with the data below. The idea is to treat the graph as a larger structure to which we have pointers, called handles (similar to Unix file handles), through which we manipulate the graph.

libhandlegraph is primarily used in vg as an abstraction layer over different backing graph implementations. It defines a common set of attributes and operations through which we can manipulate the graph. We can then use the libhandlegraph API as a layer between an underlying graph implementation and genome graph manipulation tools we plan on building.

libhandlegraph has python bindings and is now being ported to Rust. In C++ and Python, it uses the class abstraction while in Rust the trait abstraction.

libbdsg (Optimized bidirected sequence graph implementations for graph genomics) is a C++ library that provides high performance implementations of sequence graphs for graph-based pangenomics applications. Tools built on top of this are PackedGraph (low memory) and HashGraph (high-performance hash tables). vg is now using libhandlegraph through libbdsg (libhandlegraph paper).
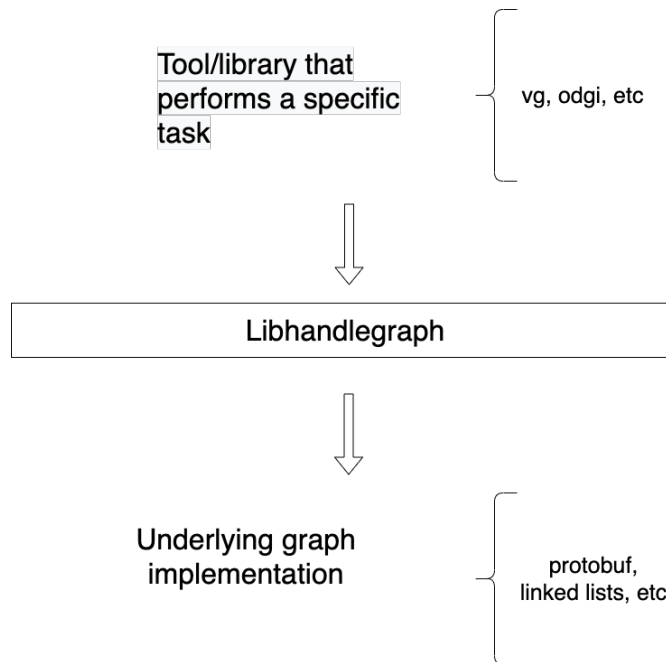
Figure 6: libhandlegraph...

7 Plaintext graphical representations

In the early 2000s assembly software was dominated by a few end to end assembly software such as SPAdes, ALLPATHS, ABySS, and SOAPdenovo https://pmelsted.wordpress.com/2014/07/17/dear-assemblers-we -need-to-talk-together/. These end to end tools made it hard to tweak parts of the assembly process which led to calls (such as THE SMALL TOOLS MANIFESTO FOR BIOINFORMATICS) for small tools that perform bits of the assembly while using plaintext files as APIs.

An early attempt was FASTG, an extension to FASTA, which is based on a directed graph (digraph) and was originally meant to represent variability in the final output of the assembly process. It encodes the sequences on arcs/edges and refers to the connection between sequences as vertices.

Like FASTA, each record contains a header line which follows the pattern a greater than sign, the edge, the neighbors of the edge and the edge properties. $> Edge : Neighbours : Properties$; where: Edge is the name given to this edge/sequence, Neighbors is a list of edges or their reverse complements that follow this edge or the reverse complement of this edge (indicated by a

preceding~), and Properties is a list of optional properties associated with this edge. To facilitate inversions, the format allows for adjacencies between forward and reverse complement. Reverse complements are indicated by a prime symbol 30.

```
>x:y;
ACGTGAGAT
```

An example of a FASTG fragment where x represents a DNA sequence and an edge in the graph. The edge is in turn followed by edge y. There exists an adjacency from edge x to edge y.

GFA (H. Li, 2016) comes in two versions: GFA1 (`https://gfa-spec.github.io/GFA-spec/GFA1.html`) and GFA2 (`https://gfa-spec.github.io/GFA-spec/GFA2.html`) with GFA2 being a superset of GFA1. Unlike FASTG, GFA is a total deviation from the FASTA format aimed specifically at plaintext representation of genome graphs and able to represent a graph at all stages of the assembly <cn> as well as varying topologies (can encode bubbles). Unlike FASTG, it encodes the sequences on the nodes, which it names segments and has edges as the connections between segments. Each line must begin with either H (header), S (Segment), F (Fragment), E (Edge), G (Gap) and G or U (Group) and each token is separated from the next by a tab (is tab delimited). It can encode extra detail through fragments which are used to specify a collection of external sequences or edges which may contain a Dazzler-trace or a CIGAR string to describe the alignment of the edge.

rGFA (H. Li et al., 2020) is GFA extended for reference (pan)genomes. It is an extension to GFA with 3 additional tags that indicate the origin of the segment to provide a unique stable coordinate system as an extension to the linear reference coordinate. Each segment is associated with one origin which forbids collapsing of different nodes from one region as would be with a cDBG in the graph by design. rGFA disallows overlaps between edges and forbids multiple edges (more than one edge between the same pair of vertices). To make use of the reference pangenome graphs GAF (H. Li et al., 2020) is a text format for sequence to graph alignment. It's an extension of PAF (H. Li, 2016). It is tab delimited like GFA.

`describe the grammar`

8   Genome graphs as databases (logic programming)

We can also treat the variation graph as a graph database. For this, SpOdgi transforms any odgi genome variation graph file into a SPARQL capable

`citation needed`

database.

9   Visualization

Visualization tools are a core tenet of bioinformatics and science in general. They help us understand our assemblies and communicate the results with others. Different tools exist depending on the level of resolution needed and the size of the graph.

GraphViz (North & Woodhull, 2002; Ellson, Gansner, Koutsofios, North, & Woodhull, 2004) is a collection of different graph visualization tools

> **expound**

Bandage (Wick, Schultz, Zobel, & Holt, 2015), originally developed for assembly graph visualization, is a standalone application written for visualizing assembly graphs. It allows the visualization of several contigs which they themselves may have various paths within them. It uses a force-directed layout via, strength is aesthetic appeal and clearly communicates components but annotation and navigation aren't possible. The major issue is the runtime scalability; force-directed layout has quadratic or even cubic costs with respect to graph size . The Open Graph Drawing Framework library (http://www.ogdf.net/) is used to perform the graph layout using the fast multipole multilevel layout algorithm, which scales well for very large graphs (Hachul and Ju nger, 2007). It reads a graph in a variety of formats: LastGraph (Velvet), FASTG (SPAdes), Trinity.fasta, ASQG and GFA and allows the export of a visualization graph either entirely or a section of it (https://rrwick.github.io/Bandage/).

> **cite pantograph docs**

MoMI-G (Yokoyama, Sakamoto, Seki, Suzuki, & Kasahara, 2019) (MOdular Multi-scale Integrated Genome graph browser) is a web based genome browser built for the visualization of structural variants (SVs) in a variation graph and has a chromosome centric view making it best for prokaryotic, <containing chromosomes> genomes. It works through a server client web architencture where the client (browser) makes requests to a backend server that one can set up locally using docker. It takes as input: a succinct representation of a variation graph in XG format, read alignment (optional), and annotations (optional).

Sequence tube maps (Beyer et al., 2019) is a javascript module that can be accessed within MoMi-G for the visualization of variation graphs or one can build their own custom API to generate the data whose aim is to represent both structural variation and sequence alignments. Tube maps were initially built to represent public transportation networks, London's iconic Tube Map, (Cartwright AM, 2012) which themselves were inspired by

circuit diagrams.

For visualizing large graphs which contain paths, assembly graphs which are de Bruijn graphs don't contain paths, it's recommended to use a pipeline such as ... These break a large graph into "chunks" that can be visualized bit by bit. Pantograph (2020) is another web based variation graph browser. It renders the genome graph in a matrix. It reads a variation graph in JSON from odgi bin.

Add image of our Household 20 dataset in pantograph

References

Baaijens, J. A., Stougie, L., & Schönhuth, A. (2020, February). Strain-aware assembly of genomes from mixed samples using flow variation graphs. *bioRxiv*, 645721. doi: 10.1101/645721

Beyer, W., Novak, A. M., Hickey, G., Chan, J., Tan, V., Paten, B., & Zerbino, D. R. (2019, December). Sequence tube maps: Making graph genomes intuitive to commuters. *Bioinformatics*, *35*(24), 5318–5320. doi: 10.1093/bioinformatics/btz597

Burrows, M., & Wheeler, D. J. (1994, May). *A Block-sorting Lossless Data Compression Algorithm* (Technical).

Cartwright AM, W. (2012). Beck's representation of London's underground system: Map or diagram? In *GSR_2* (pp. 1–14). RMIT University.

Chikhi, R., Limasset, A., & Medvedev, P. (2016, June). Compacting de Bruijn graphs from sequencing data quickly and in low memory. *Bioinformatics*, *32*(12), i201-i208. doi: 10.1093/bioinformatics/btw279

Chikhi, R., & Rizk, G. (2013, September). Space-efficient and exact de Bruijn graph representation based on a Bloom filter. *Algorithms for Molecular Biology*, *8*(1), 22. doi: 10.1186/1748-7188-8-22

Durbin, R. (2014, May). Efficient haplotype matching and storage using the positional Burrows-Wheeler transform (PBWT). *Bioinformatics*, *30*(9), 1266–1272. doi: 10.1093/bioinformatics/btu014

Ellson, J., Gansner, E. R., Koutsofios, E., North, S. C., & Woodhull, G. (2004). Graphviz and Dynagraph — Static and Dynamic Graph Drawing Tools. In M. Jünger & P. Mutzel (Eds.), *Graph Drawing Software* (pp. 127–148). Berlin, Heidelberg: Springer. doi: 10.1007/978-3-642-18638-7_6

Ferragina, P., & Manzini, G. (2000, November). Opportunistic data structures with applications. In *Proceedings 41st Annual Symposium on Foundations of Computer Science* (pp. 390–398). doi: 10.1109/SFCS.2000.892127

Flicek, P., & Birney, E. (2009, November). Sense from sequence reads: Methods for alignment and assembly. *Nature Methods*, *6*(11), S6-S12. doi: 10.1038/nmeth.1376

Garg, S., Rautiainen, M., Novak, A. M., Garrison, E., Durbin, R., & Marschall, T. (2018, July). A graph-based approach to diploid genome assembly. *Bioinformatics*, *34*(13), i105-i114. doi: 10.1093/bioinformatics/bty279

Garrison, E. (2019). *Graphical pangenomics* (Thesis, University of Cambridge). doi: 10.17863/CAM.41621

Garrison, E., Sirén, J., Novak, A. M., Hickey, G., Eizenga, J. M., Dawson, E. T., . . . Durbin, R. (2018, October). Variation graph toolkit improves read mapping by representing genetic variation in the reference. *Nature Biotechnology*, *36*(9), 875–879. doi: 10.1038/nbt.4227

Grabowski, S., Deorowicz, S., & Roguski, L. u. (2015, May). Disk-based compression of data from genome sequencing. *Bioinformatics*, *31*(9), 1389–1395. doi: 10.1093/bioinformatics/btu844

Holley, G., & Melsted, P. (2019, August). Bifrost – Highly parallel construction and indexing of colored and compacted de Bruijn graphs. *bioRxiv*, 695338. doi: 10.1101/695338

Idury, R. M., & Waterman, M. S. (1995, January). A New Algorithm for DNA Sequence Assembly. *Journal of Computational Biology*, *2*(2), 291–306. doi: 10.1089/cmb.1995.2.291

Langmead, B., Trapnell, C., Pop, M., & Salzberg, S. L. (2009, March). Ultrafast and memory-efficient alignment of short DNA sequences to the human genome. *Genome Biology*, *10*(3), R25. doi: 10.1186/gb-2009-10-3-r25

Levenshtein, V. I. (1966, February). Binary codes capable of correcting deletions, insertions and reversals. *Soviet Physics Doklady*, *10*(8), 707–710.

Li, H. (2016, July). Minimap and miniasm: Fast mapping and de novo assembly for noisy long sequences. *Bioinformatics*, *32*(14), 2103–2110. doi: 10.1093/bioinformatics/btw152

Li, H., & Durbin, R. (2009, July). Fast and accurate short read alignment with Burrows–Wheeler transform. *Bioinformatics*, *25*(14), 1754–1760. doi: 10.1093/bioinformatics/btp324

Li, H., Feng, X., & Chu, C. (2020, March). The design and construction of reference pangenome graphs. *arXiv:2003.06079 [q-bio]*.

Li, R., Yu, C., Li, Y., Lam, T.-W., Yiu, S.-M., Kristiansen, K., & Wang, J. (2009, August). SOAP2: An improved ultrafast tool for short read alignment. *Bioinformatics*, *25*(15), 1966–1967. doi: 10.1093/

bioinformatics/btp336

Li, Z., Chen, Y., Mu, D., Yuan, J., Shi, Y., Zhang, H., . . . Fan, W. (2012, January). Comparison of the two major classes of assembly algorithms: Overlap–layout–consensus and de-bruijn-graph. *Briefings in Functional Genomics*, *11*(1), 25–37. doi: 10.1093/bfgp/elr035

Limasset, A., Rizk, G., Chikhi, R., & Peterlongo, P. (2017, February). Fast and scalable minimal perfect hashing for massive key sets. *arXiv:1702.03154 [cs]*.

Manber, U., & Myers, G. (1990, January). Suffix arrays: A new method for on-line string searches. In *Proceedings of the first annual ACM-SIAM symposium on Discrete algorithms* (pp. 319–327). San Francisco, California, USA: Society for Industrial and Applied Mathematics.

Myers, E. W. (2005, September). The fragment assembly string graph. *Bioinformatics*, *21*(Suppl 2), ii79-ii85. doi: 10.1093/bioinformatics/bti1114

Needleman, S. B., & Wunsch, C. D. (1970, March). A general method applicable to the search for similarities in the amino acid sequence of two proteins. *Journal of Molecular Biology*, *48*(3), 443–453. doi: 10.1016/0022-2836(70)90057-4

North, S. C., & Woodhull, G. (2002). Online Hierarchical Graph Drawing. In P. Mutzel, M. Jünger, & S. Leipert (Eds.), *Graph Drawing* (pp. 232–246). Berlin, Heidelberg: Springer. doi: 10.1007/3-540-45848-4_19

Novak, A. M., Garrison, E., & Paten, B. (2017, July). A graph extension of the positional Burrows–Wheeler transform and its applications. *Algorithms for Molecular Biology*, *12*(1), 18. doi: 10.1186/s13015-017-0109-9

Peltola, H., Söderlund, H., & Ukkonen, E. (1984, January). SEQAID: A DNA sequence assembling program based on a mathematical model. *Nucleic Acids Research*, *12*(1Part1), 307–321. doi: 10.1093/nar/12.1Part1.307

Pevzner, P. A., Tang, H., & Waterman, M. S. (2001, August). An Eulerian path approach to DNA fragment assembly. *Proceedings of the National Academy of Sciences*, *98*(17), 9748–9753. doi: 10.1073/pnas.171285098

Putze, F., Sanders, P., & Singler, J. (2010, January). Cache-, hash-, and space-efficient bloom filters. *Journal of Experimental Algorithmics*, *14*, 4:4.4–4:4.18. doi: 10.1145/1498698.1594230

Rautiainen, M., Mäkinen, V., & Marschall, T. (2019, October). Bit-parallel sequence-to-graph alignment. *Bioinformatics*, *35*(19), 3599–3607. doi: 10.1093/bioinformatics/btz162

Roberts, M., Hayes, W., Hunt, B. R., Mount, S. M., & Yorke, J. A. (2004, December). Reducing storage requirements for biological sequence comparison. *Bioinformatics*, *20*(18), 3363–3369. doi: 10.1093/bioinformatics/bth408

Sellers, P. H. (1980, December). The theory and computation of evolutionary distances: Pattern recognition. *Journal of Algorithms*, *1*(4), 359–373. doi: 10.1016/0196-6774(80)90016-4

Sirén, J., Garrison, E., Novak, A. M., Paten, B., & Durbin, R. (2020, January). Haplotype-aware graph indexes. *Bioinformatics*, *36*(2), 400–407. doi: 10.1093/bioinformatics/btz575

Smith, T. F., & Waterman, M. S. (1981, March). Identification of common molecular subsequences. *Journal of Molecular Biology*, *147*(1), 195–197. doi: 10.1016/0022-2836(81)90087-5

Wang, A., Wang, Z., Li, Z., & Li, L. M. (2018, June). BAUM: Improving genome assembly by adaptive unique mapping and local overlap-layout-consensus approach. *Bioinformatics*, *34*(12), 2019–2028. doi: 10.1093/bioinformatics/bty020

Wick, R. R., Schultz, M. B., Zobel, J., & Holt, K. E. (2015, October). Bandage: Interactive visualization of de novo genome assemblies. *Bioinformatics*, *31*(20), 3350–3352. doi: 10.1093/bioinformatics/btv383

Yokoyama, T. T., Sakamoto, Y., Seki, M., Suzuki, Y., & Kasahara, M. (2019, November). MoMI-G: Modular multi-scale integrated genome graph browser. *BMC Bioinformatics*, *20*(1), 548. doi: 10.1186/s12859-019-3145-2